

Hani Abusharkh (habusharkh)_PreviewUser Week_14_Full_Results_Analysis.docx

Document Details

Submission ID

trn:oid::2945:412855373

Submission Date

Sep 14, 2026, 4:43 PM GMT+0

Download Date

Sep 14, 2026, 5:13 PM GMT+0

File Name

Week_14_Full_Results_Analysis.docx

File Size

37.8 KB

2 Pages

462 Words

2,614 Characters

42% detected as AI

The percentage indicates the combined amount of likely AI-generated text as well as likely AI-generated text that was also likely revised by AI.

Caution: Review required.

It is essential to understand the limitations of AI detection before making decisions about a student's work. We encourage you to learn more about Turnitin's AI detection capabilities before using the tool.



4 AI-generated 42%

Likely generated or likely generated and revised by AI.

Disclaimer

Our AI writing assessment is designed to help educators identify text that might be prepared by a generative AI tool. Our AI writing assessment may not always be accurate (it may misidentify writing that is likely human generated as AI generated and likely AI generated as human generated) so it should not be used as the sole basis for adverse actions against a student. It takes further scrutiny and human judgment in conjunction with an organization's application of its specific academic policies to determine whether any academic misconduct has occurred.

Week 14 - Full Results Analysis

Interpretation of the final reconstructed cohort and robustness checks

1. Cohort

The analysis starts with 2,657 sprint rows. Thirty FUTURE rows, 14 ACTIVE rows, and three CLOSED rows with impossible timestamps are removed before label construction. One first eligible sprint per project is removed because there is no prior history. The final model frame has 2,602 sprints.

Step	Count
Raw sprint rows	2,657
Eligible closed sprints	2,610
Cold starts removed	8
Modeled sprints	2,602
Reconstructed commitments	20,603
Modeled zero-scope sprints	911
Modeled nonempty sprints	1,691

2. Main results

Late closure remains a moderate prediction problem. The main random-forest AUC is .705. Conditional spillover is easier to rank, but the transparent scope-only logistic model reaches .885 AUC, compared with .852 for random forest. The result favors parsimony: initial scope explains most of the measurable spillover signal.

Outcome / model	Test n	Positives	F1	ROC-AUC
Late closure - Random Forest	786	246	.562	.705
Late closure - XGBoost	786	246	.535	.695
Nonempty spillover - Always positive	511	482	.971	.500
Nonempty spillover - Scope-only logistic	511	482	.696	.885
Nonempty spillover - Random Forest	511	482	.875	.852

3. Uncertainty

Project-cluster intervals are wide enough that small model differences should not be overinterpreted. Late-closure random forest has a 95% AUC interval of .604-.783. Conditional spillover has .820-.922 for scope-only logistic regression and .797-.914 for random forest.

4. Temporal and cross-project stability

Forward cut points give late-closure RF AUCs from .669 to .711 and conditional-spillover RF AUCs from .846 to .856. Held-out-project performance varies much more, especially for late closure. The evidence supports local validation rather than a single global performance claim.

5. Provenance sensitivity

Static fallback is an identification assumption whenever no change history is recorded. Requiring observed Sprint history for every committed item leaves 356 test sprints and gives .875 AUC for the scope-only model and .850 for random forest. Requiring both Sprint and Story Points history leaves only 110 test sprints; AUC falls to .772 and .652 respectively.

Subset	Test n	Scope LR AUC	RF AUC
Main nonempty cohort	511	.885	.852
Observed Sprint history	356	.875	.850
Observed Sprint + Story Points history	110	.772	.652

6. Feature evidence

Held-out permutation importance shows committed issue count is the dominant conditional-spillover feature, with an average AUC drop of .220 when shuffled. For late closure, project identity and sprint length are the largest contributors.

7. Interpretation

The strongest conclusion is not that a sophisticated model predicts Scrum success. It is that reconstructed sprint-start scope carries useful information about unfinished work, and a simple scope model captures most of it. Late closure is weaker and more project-specific. The current evidence is retrospective and should not be presented as proof that exposing scores to Scrum teams improves delivery.